Health-Aware Block Scheduling for LNG Carrier Production

ISE 572 — Reinforcement Learning Meets Shipbuilding

Stephen Eacuello Steven Blum

Spring 2026

University of Rhode Island

The Problem

A $200M Scheduling Problem

An LNG carrier costs **$200M+** and takes **2–3 years** to build.

Each ship is assembled from **200 blocks**, each passing through **8 production stages**, moved by **24 transporters** and lifted by **9 cranes**.

Today, this scheduling is done **manually** — by experienced planners using spreadsheets and intuition.

Question: Can we learn a policy that schedules better than a human expert?

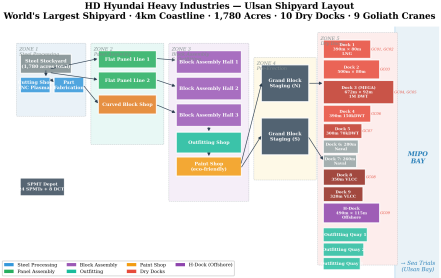

HD Hyundai Heavy Industries
Ulsan, South Korea

World's largest shipyard

What We Did

1. **Built a full simulation** of HHI Ulsan — 200 blocks, 10 docks, 9 cranes, 24 SPMTs, with NSRP-calibrated processing times
2. **Discovered entropy collapse** — PPO/SAC fail catastrophically in hierarchical masked action spaces (entropy $\rightarrow 0$ by epoch 5)
3. **Solved it with DAgger** — imitation learning achieves **118% of expert** on 50-block instances, bypassing the exploration problem entirely
4. **Exposed a scalability wall** — all methods (RL, OR, imitation) fail to scale beyond 200 blocks; simple EDD dispatch remains king at scale

How a Ship Gets Built

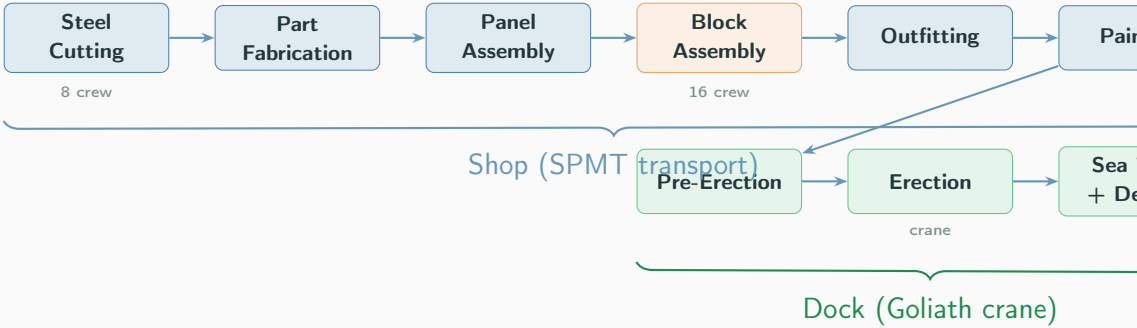

200 blocks/ship

1.24M man-hours/ship

8 production stages

Method

Formulation: Scheduling as an MDP

State: Heterogeneous graph

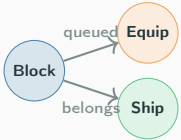

GATv2 encoder $\rightarrow$ fixed $\mathbb{R}^{256}$ embedding

Works for any number of blocks

Action: Hierarchical (10 heads)

- *What?* transport / erect / hold
- *Which SPMT?* (1 of 24)
- *Which block?* (1 of 200)
- *Which crane?* (1 of 9)

Reward: $+1$ stage advance, $+10$ ship delivery, -0.5 per tardy step

Action masking ensures only valid actions are sampled — but this creates a problem...

The Core Problem: Entropy Collapse

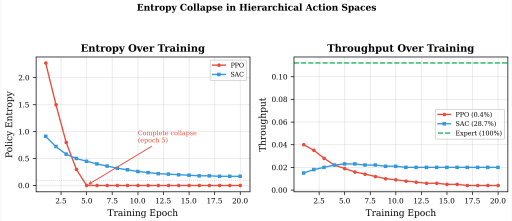

PPO entropy: **2.27** $\rightarrow$ **0.00** by epoch 5.

The policy becomes *completely deterministic* before it has learned anything useful.

Why? Action masking reduces valid options to 1–2 choices. Entropy is *mathematically bounded* by $\log k$, where k is the number of valid actions.

Result:
PPO: **0.4%** of expert

Our Solution: Learn from an Expert via DAgger

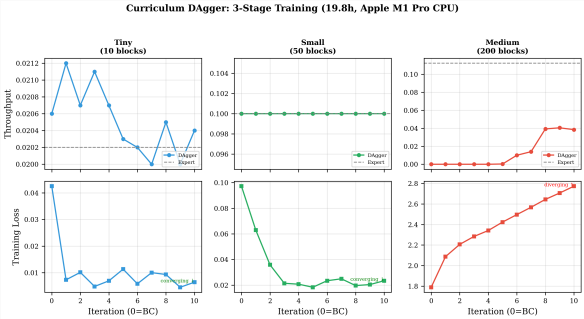

DAgger (Dataset Aggregation): Supervised learning from expert demonstrations — no exploration needed.

Curriculum: 3 stages

1. Tiny: 10 blocks
2. Small: 50 blocks
3. Medium: 200 blocks

Each stage: β anneals 0.8 $\rightarrow$ 0.2 (expert $\rightarrow$ student control)

Small instance: 118% of

Architecture

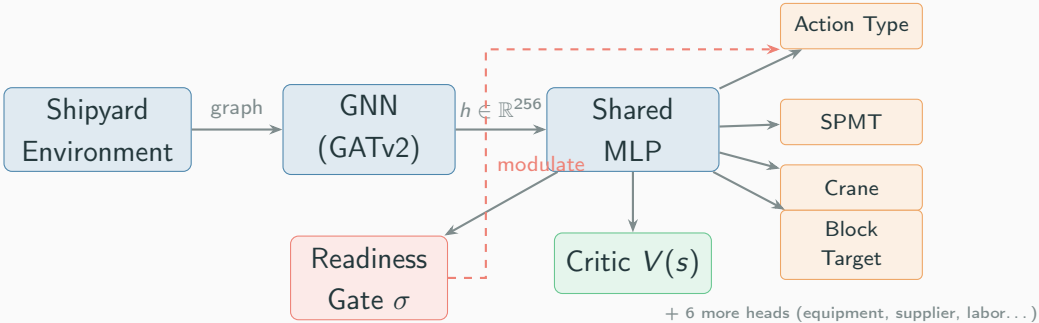

Readiness gate
logistic regression
modulates dispatch

10 action heads
only relevant heads
contribute to $\log \pi$

Action masking
invalid $\rightarrow -20$ logits
(soft, not $-\infty$)

Results

Result 1: DAgger Beats the Expert (on Small Instances)

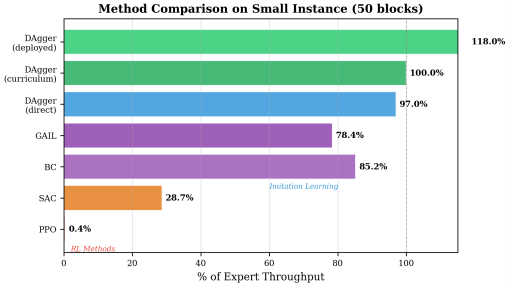

Method	vs Expert
PPO	0.4%
SAC	28.7%
BC	85.2%
GAIL	78.4%
DAgger	**118%**

DAgger *exceeds* expert by learning error recovery behaviors the rule-based EDD expert never exhibits.

10-seed validation: 97.0% $\pm$ 2.5%

Result 2: . . . But Nothing Scales to 200 Blocks

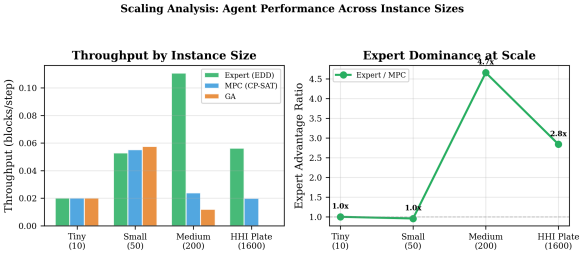

	Small (50)	Medium (200)
Expert	100%	**100%**
MPC	104%	21%
GA	109%	11%
DAgger	**118%**	0%

Expert advantage grows with scale:

1.0× (tiny) → 4.7× (medium)

GA: 1000s/episode at medium

MPC: 240K binary variables

Statistical Comparison (5 Seeds, Mann-Whitney U)

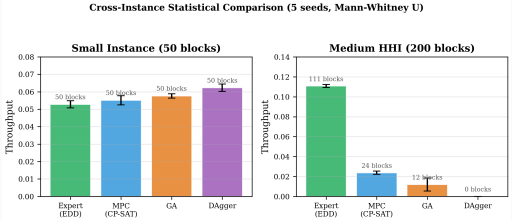

		Small	Medium
Expert	Blocks	50	111 ± 1
	Thru.	0.053	0.111
MPC	Blocks	50	24 ± 1
	Thru.	0.055	0.024
GA	Blocks	50	12 ± 5
	Thru.	0.058	0.012
DAgger	Blocks	50	0 ± 0
	Thru.	0.062	0.000

All medium pairs: $p < 0.001$

Effect sizes: $d > 25$ (huge)

Why Does DAgger Fail at Scale?

Curriculum training dynamics (medium stage):

Iter	β	Loss	Throughput
1	0.80	2.09	0.000
5	0.53	2.42	0.000
7	0.40	2.57	0.002
8	0.33	2.65	0.038
10	0.20	2.77	0.000

Loss *increases monotonically*.
The policy oscillates — learns transiently, then regresses at final checkpoint.

Root causes:

1. **Distribution shift**: 200-block state space has fundamentally different statistics than 50-block
2. **Capacity limit**: 256-dim MLP cannot represent the combinatorial 200-block action space
3. **Confirmed**: Retraining with normalizer persistence yields identical 0% result

Takeaways

Three Things to Remember

1

RL fails here

Entropy collapse in masked hierarchical action spaces is a *fundamental* problem, not a tuning issue. Validated across 3 domains.

2

Imitation works (at small scale)

DAgger achieves 118% of expert by learning error recovery. But it cannot generalize beyond its training distribution.

3

Simple rules scale best

EDD dispatch — a 1-line priority rule — outperforms MPC, GA, and learned policies at 200+ blocks. Scalability is the bottleneck.

Connection to ISE 572: Classic scheduling theory (EDD optimality for $1||\sum T_j$)

Limitations & Future Directions

What we couldn't solve

- 118% $\rightarrow$ 0% scalability cliff
- Entropy collapse (no RL fix found)
- Block assembly calibration ($R^2 = -0.03$)
- Simplified crane physics
- No weather or disruptions

What might work next

- **Attention-based heads** for variable entity counts
  (implemented, not yet trained)
- Larger policy networks (512+)
- Per-head adaptive entropy
- Transfer learning across ship types
- Digital twin integration

Demo

Live Demo: MES Dashboard

Live dashboard demo at `localhost:8050`

11 tabs: Overview · Blocks · Ships · Fleet · Health

Operations · Dependencies · GNN · Plates · Map · Docks

```
PYTHONPATH=src python experiments/live_simulation.py --policy expert
python -m src.mes.app
```

Questions?

Tests	157 passing
Instances	4 sizes (10–1600 blocks)
Agents	6 methods compared
Paper	36 pages
Dashboard	11 tabs, real-time
Training	50+ hours of experiments